

Image Alignment

CS 413 Special Topics in Computer Science

Introduction to Computer Vision







Spring 2014

Mohamed E. Hussein

Slide Sources


- Kristen Grauman's slides, Univ. of Texas at Austin

Image Alignment: Medical Image Registration



Alignment and Registration are used interchangeably to mean the same thing: mapping one image coordinate system to to another so that common pixels coincide.

Image Alignment: Mosaics



Alignment problem

- In alignment, we will fit the parameters of some **transformation** according to a set of matching feature pairs (“correspondences”).






Image alignment



- Two broad approaches:
 - Direct (pixel-based) alignment
 - Search for alignment where most pixels agree
 - Feature-based alignment
 - Search for alignment where *extracted features* agree
 - Can be verified using pixel-based alignment

Fitting an affine transformation


- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix}$$

$$\begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix}$$


Fitting an affine transformation

$$\begin{bmatrix} & & \dots & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ & & \dots & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

- How many matches (correspondence pairs) do we need to solve for the transformation parameters?
- Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for (x_{new}, y_{new}) ?
- Where do the matches come from?
 - Matching local features

Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} \dots & \dots & \dots & \dots & \dots & \dots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Fitting using homogeneous coordinates

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- The set of equations becomes

$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_i & y_i & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Recall: 2D Projective Transformations

A 2D projective transformation is also known as ***Homography***.


$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Note that, generally, $w' \neq w$ in a homography.

Projective transformations are combinations of:

- Affine transformations, and
- Projective warps

Parallel lines do not necessarily remain parallel



Recall: 2D Projective Transformations

- Since multiplying by a scalar does not change the point in homogeneous coordinates, a homography matrix can be arbitrarily scaled without changing the transformation.
- The typical ‘canonical’ form of a homography matrix is obtained by scaling the matrix such that the bottom-right most element, i , becomes 1.

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Here, the homography matrix is put in the ‘canonical’ form.

Fitting 2D Projective Transformations

To be able to use pixel coordinates of the two images directly, the equation is reformulated as:

$$w' \begin{bmatrix} x^* \\ y^* \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}, \text{ where } x^* = \frac{x'}{w'}, y^* = \frac{y'}{w'}.$$

- The set of equations becomes

$$\begin{bmatrix} x & y & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & x & y & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & x & y \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \end{bmatrix} = \begin{bmatrix} w'x^* \\ w'y^* \\ w' - 1 \end{bmatrix}$$

- Each matching pair, p_i and p'_i , gives three equations, but, adds a new unknown, w'_i .

Fitting 2D Projective Transformations

- Note that the w'_i unknowns are in the rhs, so equations need to be re-written to get all unknowns to the lhs

$$\begin{bmatrix} x & y & 1 & 0 & 0 & 0 & 0 & 0 & -x^* \\ 0 & 0 & 0 & x & y & 1 & 0 & 0 & -y^* \\ 0 & 0 & 0 & 0 & 0 & 0 & x & y & -1 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \\ w' \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ -1 \end{bmatrix}$$

- It is much easier to get rid of the w' variables all together.

Fitting 2D Projective Transformations

- For two vectors \vec{a} and \vec{b} , if $\vec{a} = \vec{b}$, then $\vec{a} \times \vec{b} = \underline{0}$, where $\underline{0}$ is a vector of all zeros.

$$\vec{a} \times \vec{b} = [\vec{a}]_{\times} \vec{b} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix} \vec{b}$$

- Back to the homography equation: let p^* and p be a pair of matching points, and let H be the homography matrix, then

$$w'p^* = Hp \implies w'p^* \times Hp = \underline{0} \implies p^* \times Hp = \underline{0}$$

- Now, we have successfully removed the w' variable
- We can get only two equation per pair of points
 - Why?
 - How many pair of points to we need to estimate?
 - Construct the resulting system of equations

An aside: Least Squares Example

Say we have a set of data points (X_1, X_1') , (X_2, X_2') , (X_3, X_3') , etc. (e.g. person's height vs. weight)

We want a nice compact formula (a line) to predict X' s from X s: $aX + b = X'$

We want to find a and b

How many (X, X') pairs do we need?

Each point gives us one equation, so, we need two points to solve for two unknowns

$$aX_1 + b = X_1'$$

$$aX_2 + b = X_2'$$

$$\begin{bmatrix} X_1 & 1 \\ X_2 & 1 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} X_1' \\ X_2' \end{bmatrix}$$

$$Ax = b$$

An aside: Least Squares Example

What if the data is noisy?


Use more points than needed and minimize fitting error

$$\begin{bmatrix} X_1 & 1 \\ X_2 & 1 \\ X_3 & 1 \\ \dots & \dots \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} X'_1 \\ X'_2 \\ X'_3 \\ \dots \end{bmatrix}$$

Overconstrained
system of
equations

Find x that achieves

$$\min \|Ax - b\|^2$$



This is called the Least
Squares solution

Recall that


$$Ax = b \Rightarrow A^T Ax = A^T b \Rightarrow x = (A^T A)^{-1} A^T b$$

It is proven that this is the least squares
solution for $Ax = b$.


Outliers

Outliers can hurt the quality of our parameter estimates, e.g.,


- an erroneous pair of matching points from two images
- an edge point that is noise, or doesn't belong to the line we are fitting.



Outliers affect least squares fit



Outliers affect least squares fit



RANSAC

RANdom Sample Consensus

Approach: we want to avoid the impact of outliers, so let's look for “inliers”, and use those only.

Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.


RANSAC: General form

RANSAC loop:


1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

RANSAC for line fitting example

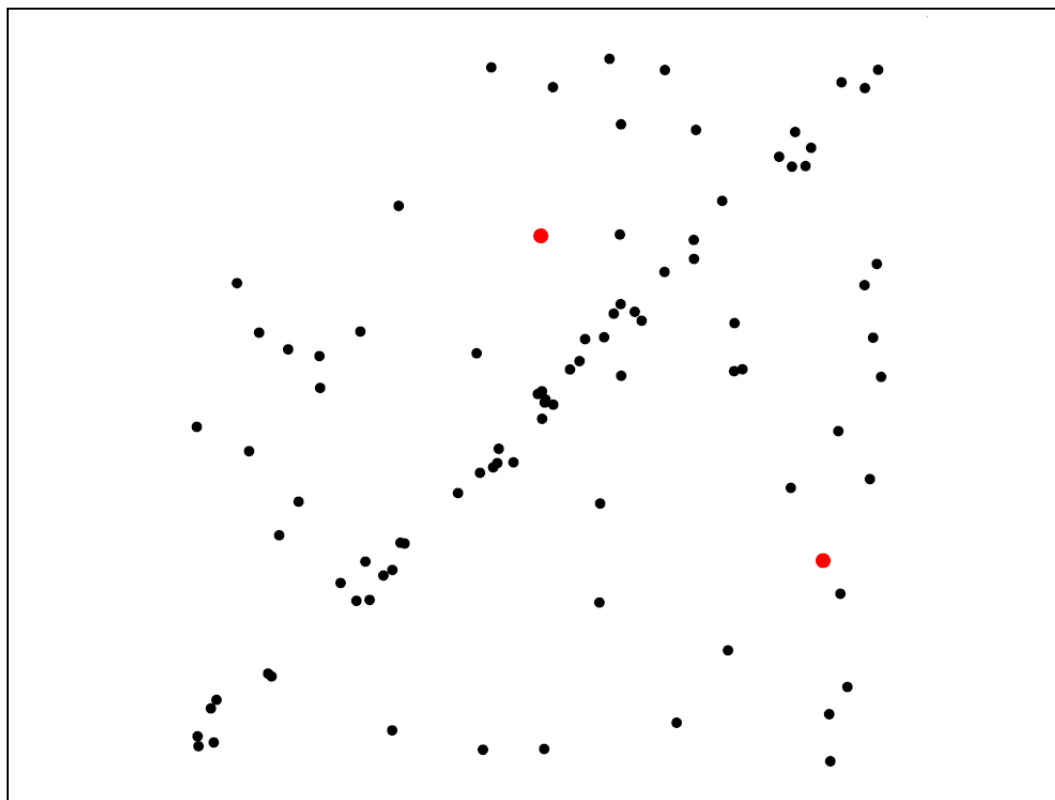


RANSAC for line fitting example



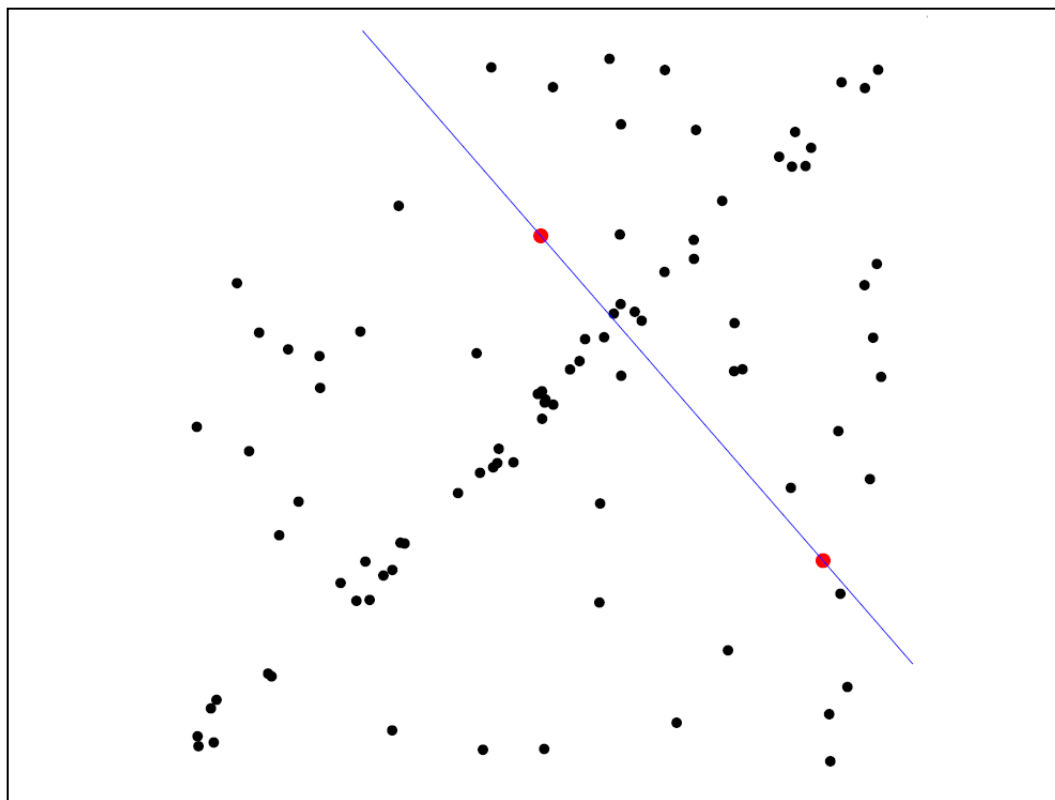
Least-squares fit

RANSAC for line fitting example




1. Randomly select minimal subset of points

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

RANSAC for line fitting example



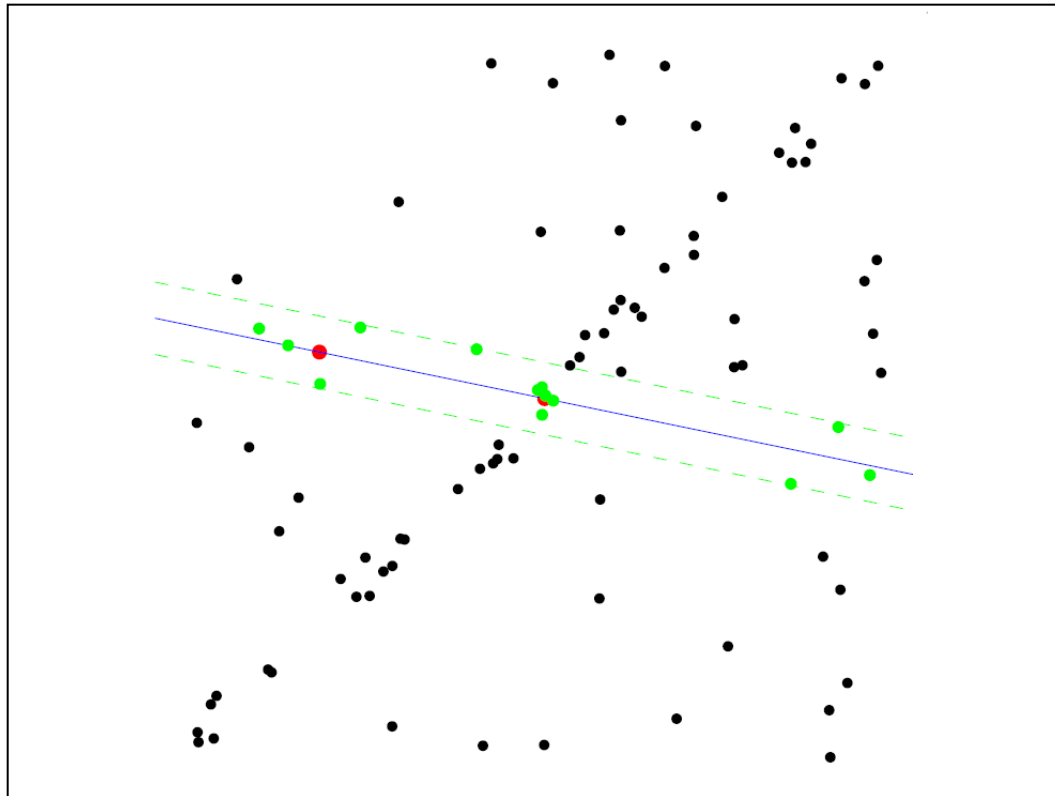
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. **Select points consistent with model**

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting

Repeat **N** times:

- Draw **s** points uniformly at random
- Fit line to these **s** points
- Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than **t**)
- If there are **d** or more inliers, accept the line and refit using all inliers

RANSAC pros and cons


- Pros
 - Simple and general
 - Applicable to many different problems
 - Often works well in practice
- Cons
 - Lots of parameters to tune
 - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
 - Can't always get a good initialization of the model based on the minimum number of samples
 - Hard to work with multiple models
 - E.g. multiple lines in an image




That is an example fitting a model
(line)...

What about fitting a transformation (translation)?


RANSAC example: Translation



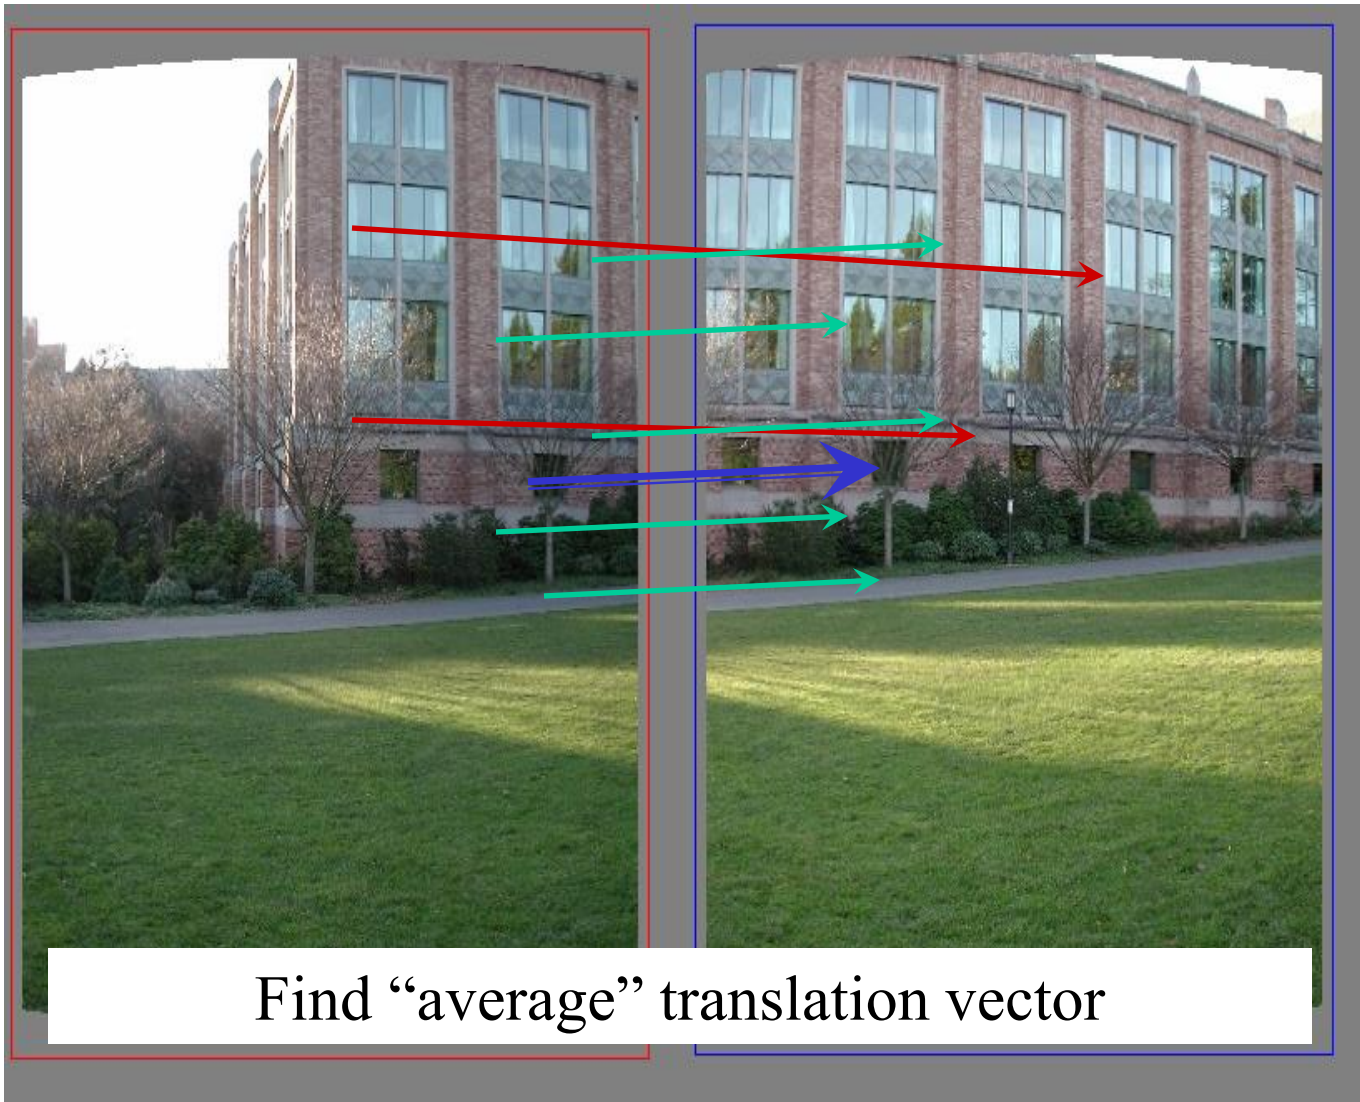
RANSAC example: Translation



RANSAC example: Translation



RANSAC example: Translation



Readings

- “Digital Image Processing”, Burger and Burge: 16.1.1-4
- “Computer Vision: Algorithms and Applications”, Richard Szeliski: 6.1.4